

A large, light gray watermark of the Stanford University seal is centered in the background. The seal is circular with a diamond-patterned border. Inside the border, the text "LELAND STANFORD JUNIOR UNIVERSITY" is written in a circular path. Below this, the German phrase "DIE LUFT DER FREIHEIT WEHT" is written. In the center of the seal is a redwood tree standing on a rocky outcrop. At the bottom of the seal, the year "1891" is inscribed.

CME 213

SPRING 2012-2013

Eric Darve

A large, light gray watermark of the Stanford University seal is centered in the background. The seal is circular with a diamond-patterned border. Inside the border, the text "LELAND STANFORD JUNIOR UNIVERSITY" is written in a circular path. Below this, the German phrase "DIE LUFT DER FREIHEIT WEHT" is written. In the center of the seal is a detailed illustration of a redwood tree standing on a rocky outcrop. At the bottom of the seal, the year "1891" is inscribed.

PTHREADS

PROGRAMMING USING THREADS

- This is the most basic approach for programming in parallel.
- Pthreads: POSIX threads. This is a standard to implement threads on UNIX systems. It is based on the C programming language.
- The other approach is based on OpenMP.

THE BASICS

- Include the header file:

`<pthread.h>`

- Compile using:

`gcc -o pth_hello pth_hello.c -lpthread`

LET'S DIVE IN

```
...
#include <pthread.h>
#define NUM_THREADS    5

void *PrintHello(void *threadid){
    long tid = (long)threadid;
    printf("Hello World! It's me, thread #%ld!\n", tid);
    pthread_exit(NULL);
}

int main(int argc, char *argv[]){
    pthread_t threads[NUM_THREADS];
    int rc; long t;
    for(t=0;t<NUM_THREADS;t++){
        printf("In main: creating thread %ld\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)t);
        if (rc){
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }

    /* Last thing that main() should do */
    pthread_exit(NULL);
}
```

OUTPUT

```
In main: creating thread 0
In main: creating thread 1
Hello World! It's me, thread #0!
In main: creating thread 2
Hello World! It's me, thread #1!
Hello World! It's me, thread #2!
In main: creating thread 3
In main: creating thread 4
Hello World! It's me, thread #3!
Hello World! It's me, thread #4!
```

THREAD CREATION

```
int pthread_create(  
    pthread_t *thread,  
    const pthread_attr_t *attr,  
    void *(*routine)(void*),  
    void *arg)
```

- **thread** thread identifier
- **routine** function that will be executed by the generated thread
- **arg** pointer to the argument value with which the thread function **routine()** will be executed
- **attr** use **NULL** for the time being

THREAD TERMINATION

A thread terminates when:

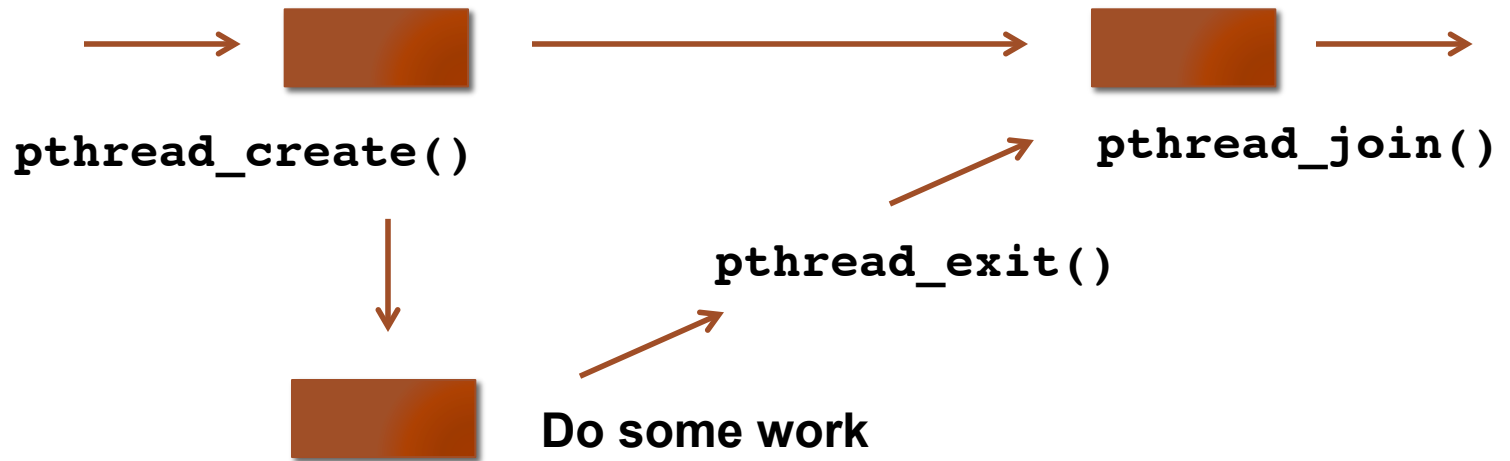
1. Thread reaches the end of its thread function, i.e., returns.
2. Thread calls

```
void pthread_exit(void *valuep)
```

Note:

- Upon termination, a thread releases its runtime stack.
- Therefore the pointer should point to: 1) a global variable, or 2) a dynamically allocated variable.

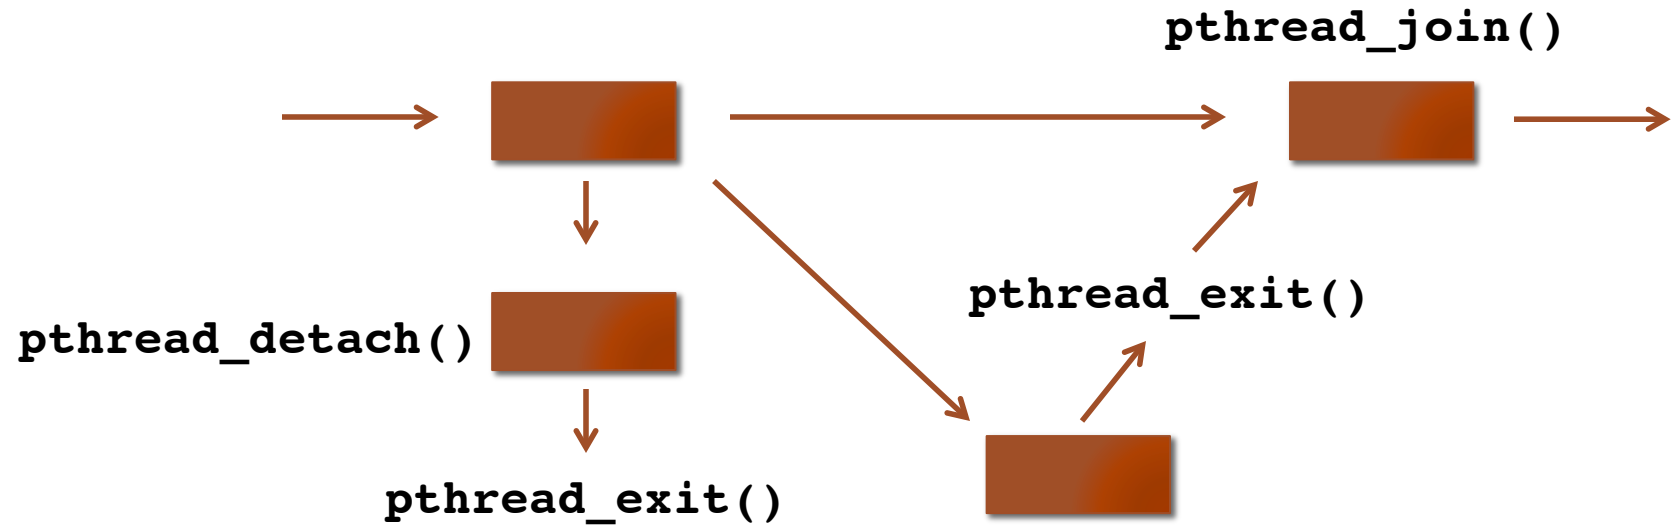
JOINING THREADS



```
int pthread_join( pthread_t thread, void **valuep )
```

- Calling thread waits for thread to terminate.
- `pthread_join` is used to synchronize threads.
- `valuep` memory address where the return value of thread is stored.

DETACHING A THREAD



DETACHING THREADS

- A thread that terminates needs to hold some internal data structure until `pthread_join()` is called.
- What happens if `pthread_join()` is never called?
 - Basically, synchronization is not needed. We launched a thread; it does its work but we don't need to know when it terminates.
- In that case, the internal data structure is never released. This can be a problem in programs that create and terminate a lot of threads.

Solution:

```
int pthread_detach(pthread_t thread)
```

- This function notifies the runtime system that the internal data structure of this thread can be freed as soon as the thread terminates.
- `pthread_join()` returns an error in that case.

EXAMPLE: MATRIX MATRIX PRODUCT

$$\mathbf{MC} = \mathbf{MA} * \mathbf{MB}$$

```

#include <pthread.h>

typedef struct {
    int size, row, column;
    double (*MA)[8], (*MB)[8], (*MC)[8];
} matrix_type_t;

void *thread_mult (void *w) {
    matrix_type_t *work = (matrix_type_t *) w;
    int i, row = work->row, column = work->column;
    work->MC[row][column] = 0;
    for (i=0; i < work->size; i++)
        work->MC[row][column] += work->MA[row][i] * work->MB[i][column];
    return NULL;
}

int main() {
    int row, column, size = 8, i;
    double MA[8][8], MB[8][8], MC[8][8];
    matrix_type_t *work;
    pthread_t thread[8*8];
    for (row=0; row<size; row++)
        for (column=0; column<size; column++) {
            work = (matrix_type_t *) malloc (sizeof (matrix_type_t));
            work->size = size;
            work->row = row;
            work->column = column;
            work->MA = MA; work->MB = MB; work->MC = MC;
            pthread_create (&(thread[column + row*8]), NULL,
                           thread_mult, (void *) work);
        }

    for (i=0; i<size*size; i++)
        pthread_join (thread[i], NULL);
}

```

The background of the slide features a large, light gray watermark of the Stanford University seal. The seal is circular and contains a redwood tree in the center, with the text 'LELAND STANFORD JUNIOR UNIVERSITY' around the top and 'DIE LUFT DER FREIHEIT WEHT' around the bottom. The year '1891' is at the very bottom. There are also stars on the sides.

THREAD COORDINATION

The background of the slide features a large, light gray watermark of the Stanford University seal. The seal is circular and contains a redwood tree in the center. The text "LELAND STANFORD JUNIOR UNIVERSITY" is written in a circle around the tree. Below the tree, the German phrase "DIE LUFT DER FREIHEIT WEHT" is inscribed. At the bottom of the seal, the year "1891" is visible. There are also several stars around the inner circle.

MUTEXES

THE RISKS OF MULTI-THREADED PROGRAMMING

- Let us assume that a well-known bank company has asked you to implement a multi-threaded code to perform bank transactions.
- You start with the modest goal of allowing deposits.
- Clients deposit money and the amount gets credited to their accounts.
- As a result of having multiple threads running concurrently the following can happen:

A PARALLEL BANK DEPOSIT

Thread 0	Thread 1	Balance
Client requests a deposit	Client requests a deposit	\$1000
Check current balance = \$1000		
	Check current balance = \$1000	
Ask for deposit amount = \$100	Ask for deposit amount = \$300	
	Compute new balance = \$1300	
Compute new balance = \$1100	Write new balance to account	\$1300
Write new balance to account		\$1100

RACE CONDITION

- Although the correct balance should be \$1,400, it is \$1,100.
- The problem is that many operations “take time” and can be “interrupted” by other threads attempting to modify the same data.
- This is called a **race condition**: the final result depends on the precise order in which the instructions are executed.
- Unless thread 0 completes its update before thread 1 (or vice versa) we get an incorrect result.
- This issue is addressed using mutexes (mutex): mutual exclusion.
- **They ensure that certain common pieces of data are accessed and modified by a single thread.**

MUTEX

- **A mutex can only be in two states: locked or unlocked.**
- **Once a thread locks a mutex:**
 - Other threads attempting to lock the same mutex are blocked
 - Only the thread that initially locked the mutex has the ability to unlock it.
- **This allows to protect regions of code.**

TYPICAL USAGE

A typical sequence in the use of a mutex is as follows:

- Create and initialize a mutex variable
- Several threads attempt to lock the mutex
- Only one succeeds and that thread owns the mutex
- The owner thread performs some set of actions
- The owner unlocks the mutex
- Another thread acquires the mutex and repeats the process

PERFORMING A DOT PRODUCT IN PARALLEL

- This is the simplest example that uses mutexes.
- All threads calculate a portion of the dot product using local variables.
- A final reduction (+) is made to a shared variable.
- This update is protected using a mutex.

HEAD OF FILE

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

typedef struct {
    double      *a;
    double      *b;
    double      sum;
    int         veclen;
} DOTDATA;

/* Define globally accessible variables and a mutex */
DOTDATA dotstr;
pthread_mutex_t mutexsum;
```

```

int main (int argc, char *argv[]) {
    int numthrds = 4;
    int veclen = 100000;
    long i;
    double *a, *b;
    pthread_t *thread;

    /* Assign storage and initialize values */
    thread = (pthread_t*) malloc(numthrds*sizeof(pthread_t));

    a = (double*) malloc (numthrds*veclen*sizeof(double));
    b = (double*) malloc (numthrds*veclen*sizeof(double));
    ...

    dotstr.vecLen = vecLen;
    dotstr.a = a; dotstr.b = b;
    dotstr.sum = 0.0;

    pthread_mutex_init(&mutexsum, NULL);

    /* Create threads to perform the dotproduct */
    for (i=0; i<numthrds; i++)
        pthread_create(&thread[i], NULL, dotprod, (void *)i);

    /* Wait on the other threads */
    for (i=0; i<numthrds; i++) pthread_join(thread[i], NULL);

    printf("Sum =  %f \n", dotstr.sum);
    free(a); free(b);
    pthread_mutex_destroy(&mutexsum);
    pthread_exit(NULL);
}

```

```

void *dotprod(void *arg) {
    int i, start, end, len;
    long offset;
    double mysum, *x, *y;
    offset = (long) arg;

    len = dotstr.vecilen;
    start = offset*len;
    end   = start + len;
    x = dotstr.a;
    y = dotstr.b;

    mysum = 0;
    for (i=start; i<end; i++) mysum += x[i] * y[i];

    /*
     Lock a mutex prior to updating the value in the shared
     structure, and unlock it upon updating.
    */
    pthread_mutex_lock(&mutexsum);
    dotstr.sum += mysum;
    printf("Thread %d did %8d to %8d: mysum = %7.3f global sum = %7.3f\n",
           (int)offset, start, end, mysum, dotstr.sum);
    pthread_mutex_unlock(&mutexsum);

    pthread_exit(NULL);
}

```


OUTPUT

```
Thread 2 did 200000 to 300000: mysum = 0.405 global sum = 0.405
Thread 1 did 100000 to 200000: mysum = 0.693 global sum = 1.099
Thread 3 did 300000 to 400000: mysum = 0.288 global sum = 1.386
Thread 0 did 0 to 100000: mysum = 12.090 global sum = 13.476
Sum = 13.476437
```

```
int pthread_mutex_init (pthread_mutex_t *mutex,  
                        const pthread_mutexattr_t *attr)
```

Initialization of mutex; choose NULL for attr.

```
int pthread_mutex_destroy (pthread_mutex_t *mutex)
```

Destruction of mutex

```
int pthread_mutex_lock (pthread_mutex_t *mutex)
```

Locks a mutex; blocks if another thread has locked this mutex and owns it.

```
int pthread_mutex_unlock (pthread_mutex_t *mutex)
```

Unlocks mutex; after unlocking, other threads get a chance to lock the mutex.